

Capstone Project Guidelines

Masters Programs

MTech / MSc in Artificial Intelligence | MSc in Business Analytics | MTech / MSc in Cybersecurity | MSc in Cloud Architecture & Security

Capstone Project Structure Overview

Each participant enrolled in the Masters program completes three capstone components as a core requirement for successful graduation. Capstone I and II are demo-ready implementations with full-length technical reports. Capstone III is a peer-reviewed publication derived from the work done in Capstone I, Capstone II, or a synthesis of both; and carries its own credit.

Three Capstones at a glance¹

Capstone I: Demo-ready implementation with full-length technical report. Industry mentor assigned.

Capstone II: Demo-ready implementation with full-length technical report. Industry mentor assigned.

Capstone III (Publication): Submission to an IEEE/Scopus-indexed or equivalent journal. Can be based on Capstone I, Capstone II, or a combined contribution. Carries independent credits. Mentor listed as second author; RACE/REVA faculty as third author.

Industry Mentor: Engaged for 30–45 days per capstone project. Minimum 5–6 structured online/offline meetings (20–30 minutes each).

Capstone Project Objectives

The Capstone Project provides a culminating, industry-connected experience that requires the integration and application of knowledge gained across the Master's program. The final certification is based on successful completion and submission of both capstone projects, with a favourable assessment from the evaluation panel during the viva.

The objective is to develop your design, implementation, and consulting capabilities by applying coursework knowledge to real-world technical challenges — and to produce research-grade work worthy of peer-reviewed publication.

Capstone projects must represent original, individual work. Describing a previous or existing project in which you were a minor contributor is not acceptable.

1. Learning Outcomes

Completion of both capstone projects will help you develop:

1. Problem-solving skills in real-world contexts: each project must derive a solution to a defined problem or opportunity in your domain.
2. Technical consulting competencies: addressing domain-specific challenges with rigour and structured methodology.

¹ Note: For the MSc Cloud Architecture and Security program, there is only one capstone project and one publication as mandatory credits.

Capstone Project Guidelines

Masters Programs

MTech / MSc in Artificial Intelligence | MSc in Business Analytics | MTech / MSc in Cybersecurity | MSc in Cloud Architecture & Security

3. Implementation proficiency: building working, demo-ready solutions using relevant tools, platforms, and frameworks.
4. Research and writing skills: producing a report of publishable quality that can be submitted to peer-reviewed venues.
5. Industry collaboration: working effectively with an assigned industry mentor to validate approach and refine outcomes.
6. Contribution to practice: advancing best practices in AI, analytics, cybersecurity, or cloud through a novel, reproducible implementation.

2. Basic Guidelines

1. Each participant² must complete three capstone components: Capstone I (implementation + report), Capstone II (implementation + report), and Capstone III (peer-reviewed publication). All three are mandatory for graduation.
2. Capstone III carries independent credits and may be derived from Capstone I, Capstone II, or a synthesis of both. The publication approach and journal selection must be discussed with your mentor and RACE faculty.
3. Capstone I and Capstone II must both be demo-ready — a working, demonstrable implementation is a mandatory deliverable alongside each report.
4. Projects must be career-enhancing and based on a live technical challenge relevant to your domain (AI, analytics, cybersecurity, or cloud).
5. Each project is an individual exercise. No group submissions.
6. Approximately 120 hours of work per project are expected to earn full credit (6/10/12 credits as per program structure).
7. Additional projects may be undertaken independently to build a personal portfolio of capstone work.
8. All reports must be pre-tested using anti-plagiarism software before submission. The University will conduct independent checks; malpractice will result in disciplinary action and report rejection.
9. Please mention REVA University and RACE affiliation in all publications and public outputs.
10. Two hardbound copies (white bond paper, white front cover with gold embossed letters) must be submitted through the Program Office. Retain a personal copy.
11. The soft copy of the project report and deck, codebase, data sources and any other artefacts must also be submitted on the LMS.
12. The project must be technology-focused, not merely a software delivery or literature exercise — the credits accrue towards an MTech/MSc with a technology concentration.

3. Industry Mentor Engagement

Each capstone project includes a structured mentoring engagement with an assigned industry professional. The mentor brings real-world context, validates your approach, and co-authors the publication where applicable. Treat mentoring sessions as professional meetings, come prepared with clear updates and specific questions.

² Note: For the MSc Cloud Architecture and Security program, there is only one capstone project and one publication as mandatory credits.

Capstone Project Guidelines

Masters Programs

MTech / MSc in Artificial Intelligence | MSc in Business Analytics | MTech / MSc in Cybersecurity | MSc in Cloud Architecture & Security

Parameter	Details
Engagement Duration	30–45 days per capstone project
Number of Sessions	5–6 structured online or offline meetings. All online sessions must be recorded.
Session Duration	20–30 minutes per session
Mode	Online (MS-Teams) or in-person as agreed
Authorship (Publication)	Mentor listed as second author; RACE/REVA faculty as third author
Mentor Assignment	Assigned by RACE program office based on domain fitment

Suggested Meeting Cadence

Session	Agenda Focus
Meeting 1	Problem scoping, domain context, and confirming project feasibility
Meeting 2	Literature review walkthrough - gaps identified, approach validated
Meeting 3	Architecture / design review - tools, datasets, and methodology sign-off
Meeting 4	Implementation check-in - progress review, troubleshooting, pivots
Meeting 5	Demo dry-run and results review - testing, validation, and benchmarking, report approval (mandatory through a mail to program office)
Meeting 6 (if needed)	Publication readiness - manuscript review, journal selection, abstract refining

Participants are responsible for scheduling sessions, preparing meeting notes, and sharing artefacts (drafts, code links, diagrams) at least 24 hours before each meeting. Mentors will provide written or verbal feedback within agreed timelines.

4. Project Report Format

The full-length report must follow the structure below. Templates for the initial pages (cover, declaration, certificate) are provided separately and must be followed strictly.

Required report sections:

1. Cover Page (format given)
2. Candidate's Declaration — signed by the candidate, guide, and co-guide
3. Acknowledgement of Project Ownership and Usage Rights
4. Company Certificate (if applicable)
5. Acknowledgements
6. Table of Contents with page references

Capstone Project Guidelines

Masters Programs

MTech / MSc in Artificial Intelligence | MSc in Business Analytics | MTech / MSc in Cybersecurity | MSc in Cloud Architecture & Security

7. List of Tables with titles and page references
8. List of Illustrations / Screenshots with titles and page references
9. List of Symbols, Abbreviations, and Nomenclature
10. Abstract
11. Chapters 1–11 (Track A: AI / Cybersecurity / Cloud) or Chapters 1–12 (Track B: Business Analytics — see Chapter Scheme, Section 5)
12. References and Bibliography
13. Appendices (if any)
14. Copies of published articles/conference papers/blogs (if any)
15. Public or shared GitHub repository with a clear README, setup instructions, and sample inputs/outputs.

5. Chapter Scheme

The report structure depends on your program track. Chapters 1–5 are common to all programs. From Chapter 6 onwards, the structure differs — follow the track that applies to you.

Which track applies to me?

Track A — AI / Cybersecurity / Cloud Architecture & Security: Chapters 1–11. Technical implementation track.

Track B — Business Analytics: Chapters 1–12. CRISP-DM methodology track.

Chapters 1–5 (Common to All Programs)

Abstract (1–2 pages)

A concise summary of the entire project: the problem, its strategic importance, the methodology, key results, and conclusions. A well-written abstract is your first impression — write it last, after all chapters are complete.

Chapter 1: Introduction (2–3 pages)

Introduce the specific problem under study. Provide context on the domain, current technical advancements, the need for this work, and the scope. State how your project contributes to organisations or communities. Figures, tables, and references are welcome here.

Chapter 2: Literature Review (2–3 pages)

A minimum of 15 references from recognised journals, technical reports, or reputable websites must support the problem formulation. The literature review must:

- Place each reviewed work in the context of its contribution to the field
- Describe relationships and contrasts between reviewed works
- Identify research gaps and unresolved contradictions

Capstone Project Guidelines

Masters Programs

MTech / MSc in Artificial Intelligence | MSc in Business Analytics | MTech / MSc in Cybersecurity | MSc in Cloud Architecture & Security

- Resolve conflicts between seemingly contradictory prior studies
- Prevent duplication by identifying what has already been done
- Point the way forward — justifying why your project matters
- Position your original contribution within the existing body of work

If your project is the basis for Capstone III (the publication), the literature review must be more extensive (20+ references) and explicitly map the gap that your project addresses — this section forms the foundation of the journal paper.

Chapter 3: Problem Statement (1–2 pages)

Describe the problem's history, setting, and organisational/community context. Identify all stakeholders, where they are located, and why they have an interest. Include the organisational structure, key processes, and resources pertinent to the problem. Discuss the strategic significance and why this problem is worthy of formal study.

Chapter 4: Objectives of the Study (2–3 pages)

Each objective must be tied to an execution phase — not just a goal, but a clear declaration of how that goal will be realised. This forces you to think about the build, not just the outcome, and creates a thread that runs from this chapter through to your conclusions.

The standard execution phases are:

- Design — specifying the architecture, model, or system structure
- Build — implementing the solution using chosen tools and datasets
- Test — measuring performance against defined metrics and a baseline
- Validate — confirming the system behaves correctly under real or simulated conditions
- Deploy — delivering a working, accessible demo or production-ready system

Some projects may add Analyse (between Test and Validate) or Optimise (between Validate and Deploy). Adapt the phases to fit your project type — but every objective must name its phase.

Example - AI Project: Intelligent Document Classifier for Insurance Claims

Objective	Statement
Objective 1: Design	Design a multi-class NLP pipeline to classify insurance claim documents into five categories (medical, motor, property, liability, life) using a transformer-based architecture.
Objective 2: Build	Fine-tune a BERT model on a labelled corpus of 12,000 insurance documents and expose the classifier as a REST API endpoint for real-time inference.
Objective 3: Test	Measure classification accuracy, precision, recall, and F1-score on a held-out test set of 2,000 documents, benchmarked against a TF-IDF + Logistic Regression baseline.

Capstone Project Guidelines

Masters Programs

MTech / MSc in Artificial Intelligence | MSc in Business Analytics | MTech / MSc in Cybersecurity | MSc in Cloud Architecture & Security

Objective	Statement
Objective 4: Validate	Validate inference latency and throughput under 100 concurrent simulated requests using Locust, targeting sub-300ms average response time.
Objective 5: Deploy	Deploy the inference API on AWS Lambda with an API Gateway front-end and demonstrate end-to-end document classification through a working web interface.

Chapter Mapping Table (mandatory element of Chapter 4)

Immediately after your objectives, include this table. Evaluators use it to navigate your report — every objective must appear here, mapped to the chapter where it is addressed.

Objective	Phase	Addressed In	What is Demonstrated
Objective 1	Design	Chapter 7	System architecture and NLP pipeline design
Objective 2	Build	Chapter 8	BERT fine-tuning, API implementation, GitHub link
Objective 3	Test	Chapter 9	Accuracy metrics, confusion matrix, baseline comparison
Objective 4	Validate	Chapter 9	Load test results, latency benchmarks
Objective 5	Deploy	Ch. 8 + Demo	Live API link, web interface walkthrough

If you cannot map an objective to a chapter, that objective needs to be revised or removed.

Example — Business Analytics Project: Customer Churn Prediction for a Telecom Operator

For Business Analytics projects using CRISP-DM, map your execution phases to the methodology stages: Understand → Prepare → Model → Evaluate → Deploy.

Objective	Statement
Objective 1: Understand	Analyse the telecom operator's business context, define churn drivers with key stakeholders, and identify the KPIs and datasets required to support a predictive churn model.
Objective 2: Prepare	Extract, clean, and transform 18 months of subscriber transaction and usage data (ETL pipeline), engineer 12 predictive features, and validate data quality against defined completeness thresholds.
Objective 3: Model	Train and compare three classification models (Logistic Regression, XGBoost, Random Forest) on an 80/20 train-test split, optimising for F1-score and AUC-ROC.
Objective 4: Evaluate	Evaluate the best-performing model against the business objective: correctly identify at least 75% of churners while maintaining precision above 0.70 to minimise unnecessary retention spend.

Capstone Project Guidelines

Masters Programs

MTech / MSc in Artificial Intelligence | MSc in Business Analytics | MTech / MSc in Cybersecurity | MSc in Cloud Architecture & Security

Objective	Statement
Objective 5: Deploy	Deploy the churn scoring model as a scheduled batch pipeline on AWS and demonstrate a live churn risk dashboard that updates weekly from new subscriber data.

Objective	Phase	Addressed In	What is Demonstrated
Objective 1	Understand	Chapter 6	Business context, KPI definition, stakeholder map
Objective 2	Prepare	Chapter 8	ETL pipeline, feature engineering, data quality report
Objective 3	Model	Chapter 9	Model comparison table, training/validation curves
Objective 4	Evaluate	Chapter 10	Business metric outcomes, precision/recall, confusion matrix
Objective 5	Deploy	Ch. 11 + Demo	Live dashboard link, batch pipeline walkthrough

Note for Business Analytics participants: Chapters 7 and 8 in your track cover Data Understanding and Data Preparation — not Software Design and Implementation. The chapter numbers in your mapping table must reference your CRISP-DM chapter scheme.

Chapter 5: Project Methodology (2 pages)

Describe the process used to execute the project. Identify all tools, platforms, and frameworks, explaining how each contributes to solving the stated problem. Document problems encountered and how they were resolved. If the methodology changed during the project, explain the changes and their rationale.

Business Analytics participants: this section must identify your chosen industry methodology. Accepted formats include CRISP-DM, TDSP, SIPOC, Q-GRID, and Cognitive Maps. For most Business Analytics capstones, CRISP-DM is the recommended choice — the chapters from 6 onwards follow this structure.

Track A: AI / Cybersecurity / Cloud Architecture & Security

Chapters 6–11

Chapter 6: Resource Requirement Specification (2–3 pages)

Outline all resources required: hardware, software, datasets, cloud infrastructure, APIs, and human resources. Justify selections. Clearly specifying requirements ensures the project is feasible, reproducible, and scalable by others.

Chapter 7: Software Design (4–5 pages)

Capstone Project Guidelines

Masters Programs

MTech / MSc in Artificial Intelligence | MSc in Business Analytics | MTech / MSc in Cybersecurity | MSc in Cloud Architecture & Security

Provide high-level design artefacts (system architecture diagram, component diagram, data flow diagram) and low-level design artefacts (state diagrams, sequence diagrams, ER diagrams, flowcharts) as appropriate to the project type. Use professional diagramming tools (draw.io, Lucidchart, PlantUML, etc.). All diagrams must be clearly labelled and referenced in the text.

Chapter 8: Implementation (4–5 pages)

Every Build and Deploy phase objective from Chapter 4 must have its own dedicated subsection here. The heading of each subsection must echo the objective it addresses. An evaluator should be able to open Chapter 4 and Chapter 8 side by side and see a clean one-to-one match.

Each subsection must include: what was built and why those specific tools/libraries/approaches were chosen, annotated pseudocode or key snippets (not full code dumps), screenshots with clear captions, and the GitHub repository link. Refer back to the objective number explicitly.

Example Subsection Structure
8.1 Implementation of Objective 2 — BERT Fine-tuning & REST API
<i>Objective: Fine-tune a BERT model on 12,000 labelled insurance documents and expose it as a REST API endpoint.</i>
[Describe the fine-tuning approach, epochs, learning rate choices, and why BERT was selected over alternatives. Include an annotated snippet of the training loop and a screenshot of training loss curves. State the GitHub link.]
8.2 Implementation of Objective 5 — Cloud Deployment & Web Interface
<i>Objective: Deploy the inference API on AWS Lambda and demonstrate end-to-end document classification.</i>
[Describe the Lambda deployment configuration, API Gateway setup, and the web interface. Include architecture screenshot and the live demo link.]

Design-phase objectives (Objective 1 in the example above) are addressed in Chapter 7, not here. Test and validate objectives are addressed in Chapter 9. Do not repeat those descriptions in Chapter 8, cross-reference them instead.

Chapter 9: Testing and Validation (3–4 pages)

Design test cases for each functional requirement. Execute and record results in text, tables, or screenshots/graphs. Validate your results against theoretical solutions, experimental benchmarks, or published reference implementations. If no reference is available, use a standard benchmark problem. Explain all trends in

Capstone Project Guidelines

Masters Programs

MTech / MSc in Artificial Intelligence | MSc in Business Analytics | MTech / MSc in Cybersecurity | MSc in Cloud Architecture & Security

your graphs and tables using a theoretical basis. The project demo must also be validated here — include the demo link or recorded demo.

Chapter 10: Analysis and Results (2–3 pages)

Connect every result back to the objective it addresses. Do not present results in isolation — for each key metric or finding, explicitly state which objective it validates and whether the target was met.

Example: 'Objective 3 targeted an F1-score above 0.85. The achieved F1 of 0.91 on the held-out test set confirms this objective was met. The baseline TF-IDF model achieved F1 of 0.73, demonstrating a 24.7% relative improvement.' Include numerical or descriptive analysis — performance metrics, comparative benchmarking, accuracy/precision/recall tables, SWOT analysis — as appropriate to your domain.

Chapter 11: Conclusions and Recommendations for Future Work (2–3 pages)

Address each objective from Chapter 4 by name and state whether it was fully met, partially met, or not met — and why. Use numbers. Do not extrapolate beyond what your results support.

Example: 'All five objectives were addressed. Objectives 1–3 were fully met. Objective 4 (latency validation) was partially met: average response time of 340ms fell short of the 300ms target under peak load, pointing to the need for model quantisation or caching in future work. Objective 5 was fully met.' Close with honest recommendations for extending the work — these often become the starting point for Capstone III or further research.

Track B: Business Analytics (CRISP-DM)

Chapters 6–12

Business Analytics capstones follow the CRISP-DM (Cross-Industry Standard Process for Data Mining) framework from Chapter 6 onwards. Chapters 1–5 remain identical to Track A. The chapter numbering is different; your Chapter Mapping Table in Chapter 4 must reference the CRISP-DM chapter numbers below.

Chapter 6: Business Understanding (2–3 pages)

Engage with key stakeholders (business heads, operations leads, end users) to understand the current business challenge in depth. The goal of this chapter is to translate a business problem into a data mining problem definition, and to produce a preliminary project plan aligned to business objectives.

This chapter must cover: the business context and strategic importance of the problem; the stakeholder landscape and their specific concerns; the definition of success from a business perspective (not just a technical one); and an articulation of the monetary or operational value — cost savings, revenue uplift, risk reduction — that a successful model would deliver.

Capstone Project Guidelines

Masters Programs

MTech / MSc in Artificial Intelligence | MSc in Business Analytics | MTech / MSc in Cybersecurity | MSc in Cloud Architecture & Security

Chapter 7: Data Understanding (4–5 pages)

Develop a clear picture of the data available to solve the problem. This chapter bridges the business understanding and the technical implementation — a weak data understanding chapter almost always produces a weak model.

This chapter must include: a description of all data sources and how they were accessed; an ER diagram or data schema showing entity relationships; preliminary metrics, KPIs, and hypotheses to be tested; and initial observations about data quality, completeness, and suitability. All datasets must be identified, with provenance (source, size, time range, access method) clearly documented.

Chapter 8: Data Preparation (4–5 pages)

Describe every step taken to transform raw data into the final, model-ready dataset. This is typically the most time-intensive phase of a data science project — document it thoroughly.

This chapter must cover the full ETL (Extract, Transform, Load) process: data extraction and ingestion; handling of missing values, duplicates, and outliers; feature engineering and selection; encoding, normalisation, and scaling decisions; and the structure of the final train/test split. Include data quality metrics before and after cleaning.

Exploratory analysis and visualisation also belong here: distribution plots, correlation heatmaps, hypothesis test results, and descriptive statistics. Each visualisation must be interpreted — do not include charts without commentary.

Each Prepare-phase objective from Chapter 4 must have its own dedicated subsection here, structured the same way as the Implementation chapter in Track A. The subsection heading must echo the objective.

Chapter 9: Modelling (3–4 pages)

Based on the hypotheses and business objectives established in Chapter 6, select and justify the modelling techniques to be applied. Document the rationale for each model choice — not just what was used, but why it is appropriate for this problem type and dataset.

This chapter must cover: the modelling techniques selected and the justification for each; hyperparameter settings and tuning approach; the train/test split strategy (and cross-validation methodology if used); and the procedure followed during model training. Each Model-phase objective from Chapter 4 must be addressed here with a dedicated subsection.

Include training and validation performance curves, feature importance plots (where applicable), and any intermediate results that informed decisions during the modelling phase.

Chapter 10: Model Evaluation (1–2 pages)

Capstone Project Guidelines

Masters Programs

MTech / MSc in Artificial Intelligence | MSc in Business Analytics | MTech / MSc in Cybersecurity | MSc in Cloud Architecture & Security

Evaluate the trained models against both technical metrics and the business objectives defined in Chapter 6. A model that achieves high accuracy but fails to meet the business objective is not a successful outcome — address both dimensions explicitly.

Connect every evaluation result back to the objective it addresses. State clearly whether the business target was met.

Example: 'Objective 4 targeted a precision above 0.70 while correctly identifying at least 75% of churners. The XGBoost model achieved precision of 0.76 and recall of 0.78 on the held-out test set, confirming Objective 4 was fully met. The Logistic Regression baseline achieved precision of 0.61 and recall of 0.68, demonstrating a 24.6% precision improvement.'

Chapter 11: Deployment (1–2 pages)

Describe the deployment of the validated model into a production or near-production environment. This chapter is the equivalent of the Deploy-phase objective — it must demonstrate that the solution works in a real or simulated operational context, not just in a notebook.

Cover: the deployment architecture (batch pipeline, real-time API, dashboard, or scheduled job); the infrastructure used (cloud platform, containerisation, scheduling); how predictions are surfaced to business users; and any monitoring or refresh mechanism. Include the live demo link or a recorded demo video. All code must be in a GitHub repository with a clear README.

Chapter 12: Analysis and Results (2–3 pages)

Present the overall outcomes of the project. Connect technical results to the business objectives defined in Chapter 6 and the objectives from Chapter 4. This chapter bridges the deployment and the conclusions — show how effectively the proposed solution meets its intended business purpose.

Include: a summary of model performance against each evaluation metric; the business value delivered or projected (cost savings, churn reduction rate, revenue impact); comparative analysis against the baseline; and a SWOT analysis of the solution in its operational context. Address each objective from Chapter 4 by name, stating whether it was fully met, partially met, or not met.

Example: 'All five objectives were addressed. Objectives 1–4 were fully met. Objective 5 (deployment) was partially met — the batch pipeline is operational but the real-time dashboard is pending stakeholder sign-off on data access permissions, identified as future work.'

6. Capstone III: Publication Requirement

Capstone III is the publication component of your Master's program and carries independent credits. It is based on the research and implementation work from Capstone I, Capstone II, or a combined contribution spanning both projects. This is a mandatory graduation requirement, not optional.

Capstone Project Guidelines

Masters Programs

MTech / MSc in Artificial Intelligence | MSc in Business Analytics | MTech / MSc in Cybersecurity | MSc in Cloud Architecture & Security

Publication Mandate

The journal manuscript must be submitted within 7 days of completing the final viva. Proof of submission (acceptance confirmation is preferred but submission proof is the minimum for graduation clearance) must be included in the final submission folder. The source project(s)- Capstone I, II, or combined, must be declared at the time of submission.

Authorship order: Student (First Author) | Industry Mentor (Second Author) | RACE/REVA Faculty (Third Author)

Accepted venues: IEEE Transactions, Elsevier (ScienceDirect), Springer, ACM, or any Scopus/UGC-CARE/Web of Science indexed journal with a verifiable Q-rating.

RACE will support the publication process through faculty co-authorship, manuscript review feedback, and guidance on journal selection. Work closely with your mentors, if the quality of the work merits it, target Q1 or Q2-rated journals in your domain.

7. Demo Readiness

Both capstone projects must produce a working, demonstrable system — not a prototype or wireframe. The demo must be stable, accessible, and structured around the project's stated objectives.

1. Provide a live demo link (hosted on cloud or accessible during Viva) or a recorded demo video in exceptional cases (3–5 minutes, approval is necessary from the program office), clearly narrated.
2. The demo must map directly to the objectives stated in Chapter 4; reviewers will check coverage.
3. Code must be in a public or shared GitHub repository with a clear README, setup instructions, and sample inputs/outputs.
4. If the project involves a model, include inference examples and performance benchmarks in the demo.
5. For cybersecurity or cloud projects, document infrastructure/architecture and include screenshots of deployed environments.

8. Final Presentation and Defence

You must defend both capstone projects during the viva. The viva panel will include a mix of industry and academic experts. Evaluation is based on:

Evaluation Parameter	Weightage
Topic Selection, Business Understanding, Literature Review & Methodology	20%
Project Implementation (code, architecture, demo readiness)	20%
Project Report (quality, completeness, writing)	30%
Presentation and Defence (clarity, depth, Q&A)	30%

Capstone Project Guidelines

Masters Programs

MTech / MSc in Artificial Intelligence | MSc in Business Analytics | MTech / MSc in Cybersecurity | MSc in Cloud Architecture & Security

Final Submission Package

Submit a zipped folder to the RACE program office with all files named and organised clearly. The folder must contain:

1. Full project report (MS Word, not more than 40–50 pages); hardbound copy to be handed to the panel before viva
2. Code and scripts in a GitHub repository (link included in the report)
3. Presentation deck (MS PowerPoint)
4. Demo link or recorded demo video
5. Proof of journal submission/acceptance (Capstone III — required for graduation clearance)
6. Anti-plagiarism report